



CoGrammar

Datasets and DataFrames



**SKILLS
FOR LIFE**

SKILLS BOOTCAMPS



Department
for Education

Data Science Lecture Housekeeping

- The use of disrespectful language is prohibited in the questions, this is a supportive, learning environment for all - please engage accordingly.
(FBV: Mutual Respect.)
- No question is daft or silly - **ask them!**
- There are **Q&A sessions** midway and at the end of the session, should you wish to ask any follow-up questions. Moderators are going to be answering questions as the session progresses as well.
- If you have any questions outside of this lecture, or that are not answered during this lecture, please do submit these for upcoming Open Classes.
You can submit these questions here: [Open Class Questions](#)

Data Science Lecture Housekeeping cont.

- For all **non-academic questions**, please submit a query:
www.hyperiondev.com/support
- Report a **safeguarding** incident:
www.hyperiondev.com/safeguardreporting
- We would love your **feedback** on lectures: [Feedback on Lectures](#)

Lecture Objectives

- **Learn how to read and manipulate data with the power of Pandas**

Working With Datasets

- ★ For context, when a dataset is mentioned, it is a collection of related data.
- ★ With datasets we can manipulate the data in a multitude of ways programmatically.
- ★ With the help of pandas DataFrames, we can effortlessly manipulate data to suit our needs.

Jupyter Notebook

- ★ From now on, we will be using Jupyter Notebook. Which can be described as follows from its official website (jupyter.org):
 - “The Jupyter Notebook is an open-source web application that allows you to create and share documents that contain live code, equations, visualisations and narrative text.”
 - “Uses include data cleaning and transformation, numerical simulation, statistical modelling, data visualisation, machine learning and much more.”

Pandas DataFrame

- ★ The pandas' library documentation defines a DataFrame as a “two-dimensional, size-mutable, with labelled rows and columns.”

columns axis=1

column name

more columns to display

index label

index axis=0

missing values

data (values)

	color	director_name	num_critic_for_reviews	duration	...	actor_2_facebook_likes	imdb_score	aspect_ratio	movie_facebook_likes
0	Color	James Cameron	723.0	178.0	...	936.0	7.9	1.78	33000
1	Color	Gore Verbinski	302.0	169.0	...	5000.0	7.1	2.35	0
2	Color	Sam Mendes	602.0	148.0	...	393.0	6.8	2.35	85000
3	Color	Christopher Nolan	813.0	164.0	...	23000.0	8.5	2.35	164000
4	NaN	Doug Walker	NaN	NaN	...	12.0	7.1	NaN	0

Anatomy of a DataFrame

Pandas DataFrame

- ★ To simplify, think of a DataFrame as a table of data with the following characteristics (Lynn 2018) :
 - “There can be multiple rows and columns in the data
 - Each row represents a sample of data,
 - Each column contains a different variable that describes the samples (rows).
 - The data in every column is usually the same type of data - eg. numbers, strings, dates.
 - Usually, unlike an excel dataset, DataFrames avoid having missing values.

Working With Datasets

- ★ You can read data from a .csv (comma separated values) file into a DataFrame using the `read_csv()` function.

```
pd.read_csv('credit.csv', delimiter = ',')
```

- ★ There are other functions that can be used to read data from other sources into DataFrames, such as : `read_excel()`, `read_sql()` to name a few.

Selecting Columns in Pandas

- ★ There are a multitude of ways to specify columns in Pandas. A simplified way would be to use dictionary notation for specific columns.
- ★ We could think of DataFrames can be thought of as dictionaries : keys would be the column name and values would be the values within the column.

Example

```
import pandas as pd
import seaborn as sns

df = sns.load_dataset('iris')
print(df.columns)

# pulling data from the species column
species = df['species']
print(species)

# If you want multiple columns we can feed a list of columns we need
multi = df[['species', 'petal_length', 'sepal_length']]
print(multi)
```

Built-in DataFrame Methods

- ★ Here is a list of common built-in functions in Pandas for such things :
- **mean()** - Computes the mean for each column.
 - **min()** - Computes the minimum for each column.
 - **max()** - Computes the maximum for each column.
 - **std()** - Computes the standard deviation for each column.
 - **var()** - Computes the variance for each column.
 - **nunique()** - Computes the number of unique values in each column.

Grouping in Pandas

- ★ Data analysis can get a bit complicated at times, and some more advanced functionality might be needed.
- ★ For instance, we need the average the insurance charges for all individuals between the age of 30 to 35.

Example

```
import pandas as pd

df = pd.read_csv('insurance.csv')

below_35 = df[df['age'] < 35]
between_30_and_35 = below_35[below_35['age'] < 30]

print(between_30_and_35['charges'].mean())
```

Working With Datasets

- ★ While the previous example works nicely, what if we wanted to average the charges for every age group?
- ★ It can be done with the same syntax, but will take quite a few lines of code.
- ★ Luckily, pandas has something for us to make this possible using the least lines of code possible :

```
all_ages = df.groupby('age')['charges'].mean()  
  
print(all_ages)
```

CoGrammar

Q & A SECTION

**Please use this time to ask
any questions relating to the
topic, should you have any.**



CoGrammar

Thank you for joining!